

## train\_model\_fixed\_lr

May 27, 2024

```
[1]: import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.utils.data import DataLoader
from bird_song_dataset import BirdSongDataset, DataPaths, DeviceManager
from torch.utils.tensorboard import SummaryWriter
import os
```

```
[2]: class SimpleCNN(nn.Module):
    """
    CNN model for image classification

    This network consists of two convolutional layers followed by two fully
    ↪connected layers
    The network uses ReLU activation functions for non-linearity and max
    ↪pooling for down-sampling

    Nueral net architecture:
    - conv1:
        The first convolutional layer holds 16 filters, a kernel size of 3,
    ↪stride of 1, and padding of 1
    - conv2:
        The second convolutional layer holds 32 filters, a kernel size of
    ↪3, stride of 1, and padding of 1
    - fc1:
        The first fully connected layer that maps from the flattened output
    ↪of the last pooling layer to 512 features
    - fc2:
        The second fully connected layer that maps the 512 features to the
    ↪number of classes

    The forward method defines the data flow through the network, applying
    ↪layers sequentially with ReLU activation functions and pooling operations
    """

    def __init__(self, num_classes=5):
```

```

super(SimpleCNN, self).__init__()
# First convolutional layer with 16 filters
self.conv1 = nn.Conv2d(1, 16, kernel_size=3, stride=1, padding=1)

# Second convolutional layer with 32 filters
self.conv2 = nn.Conv2d(16, 32, kernel_size=3, stride=1, padding=1)

# First fully connected layer, transforming the feature map from
↳ convolutional layers into a 512-dimensional vector
self.fc1 = nn.Linear(32768, 512)

# Final fully connected layer that outputs probability distribution
↳ across the classes
self.fc2 = nn.Linear(512, num_classes)

def forward(self, x):
    # Apply the first convolutional layer followed by ReLU activation and
    ↳ max pooling
    x = F.relu(self.conv1(x))
    x = F.max_pool2d(x, 2)

    # Apply the second convolutional layer followed by ReLU activation and
    ↳ another max pooling
    x = F.relu(self.conv2(x))
    x = F.max_pool2d(x, 2)

    # Flatten the output from the convolutional layers to prepare for the
    ↳ fully connected layer
    x = torch.flatten(x, 1)

    # Apply the first fully connected layer with ReLU activation
    x = F.relu(self.fc1(x))

    # Output layer that maps to the number of classes
    x = self.fc2(x)

    return x

```

```

[3]: # Get dynamic paths
data_paths = DataPaths()
paths = data_paths.get_paths()
print(paths.keys())

```

```

dict_keys(['csv_file_path', 'wav_files_dir', 'models_dir', 'results_dir',
'runs_dir'])

```

```
[4]: # Instantiate dataset class
bird_dataset = BirdSongDataset(csv_file=paths['csv_file_path'],
    ↪root_dir=paths['wav_files_dir'])
print(f"Dataset size: {len(bird_dataset)}")
```

Dataset size: 5422

```
[5]: # Data split sizes for train, val, and test
train_size = int(0.7 * len(bird_dataset))
val_size = int(0.15 * len(bird_dataset))
test_size = len(bird_dataset) - train_size - val_size

print(f'Data split sizes for train, val, and test: {train_size, val_size,
    ↪test_size}')
```

Data split sizes for train, val, and test: (3795, 813, 814)

```
[6]: # Random split
train_dataset, val_dataset, test_dataset = torch.utils.data.
    ↪random_split(bird_dataset, [train_size, val_size, test_size])

train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=64, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)

# Determine accelerator device
device_manager = DeviceManager()
device = device_manager.device
```

Using MPS (Apple Silicon GPU)

```
[7]: # Define the model, loss function, and optimizer
model = SimpleCNN(num_classes=5).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Initialize TensorBoard writer for logging
writer = SummaryWriter(f"{paths['runs_dir']}/bird_song_experiment_fixed_lr")

# Initialize the best validation loss to a high value
best_val_loss = float('inf')

# Set number of epochs for training
num_epochs = 10
for epoch in range(num_epochs):
    # Set model to training mode and initialize running loss
    model.train()
```

```

running_loss = 0.0

# Loop over batches in the training dataset
for batch in train_loader:
    inputs, labels = batch['spectrogram'].to(device), batch['label'].
    ↪to(device)
    # Zero the gradients
    optimizer.zero_grad()
    # Forward pass
    outputs = model(inputs)
    # Compute loss
    loss = criterion(outputs, labels)
    # Backward pass
    loss.backward()
    # Update parameters
    optimizer.step()
    # Accumulate the loss
    running_loss += loss.item()

# Compute and log training loss
training_loss = running_loss / len(train_loader)
writer.add_scalar('Loss/train', training_loss, epoch)

# Set model to evaluation mode and compute validation loss
model.eval()
val_loss = 0.0
with torch.no_grad():
    # Loop over batches in the validation dataset
    for batch in val_loader:
        # Extract inputs and labels from the batch
        inputs, labels = batch['spectrogram'].to(device), batch['label'].
        ↪to(device)
        # Forward pass: compute model output
        outputs = model(inputs)
        # Compute loss
        loss = criterion(outputs, labels)
        # Accumulate the validation loss over all of the batches
        val_loss += loss.item()

# Logging the validation loss
validation_loss = val_loss / len(val_loader)
writer.add_scalar('Loss/validation', validation_loss, epoch)

# Check if this is the best model so far
if validation_loss < best_val_loss:
    best_val_loss = validation_loss
    # Save the model

```

```

        torch.save(model.state_dict(), f"{paths['models_dir']}/model_fixed_lr/
↪model_best.pth")

        print(f"Epoch {epoch+1}/{num_epochs}, Training Loss: {training_loss},
↪Validation Loss: {validation_loss}")
        print("-" * 75)

# Closing TensorBoard writer
writer.close()

```

Epoch 1/10, Training Loss: 29.40092813372612, Validation Loss:  
1.0550424869243915

-----  
Epoch 2/10, Training Loss: 0.8446318248907725, Validation Loss:  
0.8775509870969332

-----  
Epoch 3/10, Training Loss: 0.5641969790061315, Validation Loss:  
0.8226675391197205

-----  
Epoch 4/10, Training Loss: 0.3868511237204075, Validation Loss:  
0.8405184975037208

-----  
Epoch 5/10, Training Loss: 0.2856514650086562, Validation Loss:  
0.8824616991556608

-----  
Epoch 6/10, Training Loss: 0.22764844683309396, Validation Loss:  
0.8985623854857224

-----  
Epoch 7/10, Training Loss: 0.14086748684446018, Validation Loss:  
0.9348215598326463

-----  
Epoch 8/10, Training Loss: 0.07668051312988003, Validation Loss:  
1.011623501777649

-----  
Epoch 9/10, Training Loss: 0.06212049775446455, Validation Loss:  
1.0507894937808697

-----  
Epoch 10/10, Training Loss: 0.03388605825603008, Validation Loss:  
1.2411040709568903