

train_model_sch_lr_es

May 27, 2024

```
[1]: import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.utils.data import DataLoader, random_split
from torch.utils.tensorboard import SummaryWriter
from bird_song_dataset import BirdSongDataset, DataPaths, DeviceManager
from torchvision.transforms import ToTensor
from torch.optim.lr_scheduler import StepLR
from sklearn.model_selection import train_test_split
from datetime import datetime
```

```
[2]: class SimpleCNN(nn.Module):
    """
    CNN model for image classification

    This network consists of two convolutional layers followed by two fully
    ↪connected layers
    The network uses ReLU activation functions for non-linearity and max
    ↪pooling for down-sampling

    Nueral net architecture:
    - conv1:
        The first convolutional layer holds 16 filters, a kernel size of 3,
    ↪stride of 1, and padding of 1
    - conv2:
        The second convolutional layer holds 32 filters, a kernel size of
    ↪3, stride of 1, and padding of 1
    - fc1:
        The first fully connected layer that maps from the flattened output
    ↪of the last pooling layer to 512 features
    - fc2:
        The second fully connected layer that maps the 512 features to the
    ↪number of classes

    The forward method defines the data flow through the network, applying
    ↪layers sequentially with ReLU activation functions and pooling operations
```

```

"""

def __init__(self, num_classes=5):
    super(SimpleCNN, self).__init__()
    # First convolutional layer with 16 filters
    self.conv1 = nn.Conv2d(1, 16, kernel_size=3, stride=1, padding=1)
    # Second convolutional layer with 32 filters
    self.conv2 = nn.Conv2d(16, 32, kernel_size=3, stride=1, padding=1)
    # First fully connected layer, transforming the feature map from
    ↪ convolutional layers into a 512-dimensional vector
    self.fc1 = nn.Linear(32768, 512)
    # Final fully connected layer that outputs probability distribution
    ↪ across the classes
    self.fc2 = nn.Linear(512, num_classes)

    def forward(self, x):
        # Apply the first convolutional layer followed by ReLU activation and
    ↪ max pooling
        x = F.relu(self.conv1(x))
        x = F.max_pool2d(x, 2)
        # Apply the second convolutional layer followed by ReLU activation and
    ↪ another max pooling
        x = F.relu(self.conv2(x))
        x = F.max_pool2d(x, 2)
        # Flatten the output from the convolutional layers to prepare for the
    ↪ fully connected layer
        x = torch.flatten(x, 1)
        # Apply the first fully connected layer with ReLU activation
        x = F.relu(self.fc1(x))
        # Output layer that maps to the number of classes
        x = self.fc2(x)

    return x

```

```

[3]: # Get dynamic paths
data_paths = DataPaths()
paths = data_paths.get_paths()
print(paths.keys())

```

```

dict_keys(['csv_file_path', 'wav_files_dir', 'models_dir', 'results_dir',
'runs_dir'])

```

```

[4]: # Instantiate dataset class
bird_dataset = BirdSongDataset(csv_file=paths['csv_file_path'],
    ↪ root_dir=paths['wav_files_dir'])
print(f"Dataset size: {len(bird_dataset)}")

```

Dataset size: 5422

```
[5]: # Data split sizes for train, val, and test
train_size = int(0.7 * len(bird_dataset))
val_size = int(0.15 * len(bird_dataset))
test_size = len(bird_dataset) - train_size - val_size

print(f'Data split sizes for train, val, and test: {train_size, val_size,
↪test_size}')
```

Data split sizes for train, val, and test: (3795, 813, 814)

```
[6]: # Get the labels from the dataset for stratification
labels = bird_dataset.labels
labels
```

```
[6]: array([1, 1, 1, ..., 2, 2, 2])
```

```
[7]: # Random split
train_dataset, val_dataset, test_dataset = random_split(bird_dataset,
↪[train_size, val_size, test_size])

train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=64, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)

# Determine accelerator device
device_manager = DeviceManager()
device = device_manager.device
print(device)
```

Using MPS (Apple Silicon GPU)

mps

```
[8]: # Define the model, loss function, optimizer, and learning rate scheduler
model = SimpleCNN(num_classes=5).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
scheduler = StepLR(optimizer, step_size=3, gamma=0.1)

# Initialize TensorBoard writer for logging
writer = SummaryWriter(f"{paths['runs_dir']}/
↪bird_song_experiment_with_scheduler_early_stopping")

# Initialize variables for early stopping mechanism
patience = 3
best_val_loss = float('inf')
epochs_no_improve = 0
```

```

early_stop = False

# Set number of epochs for training
num_epochs = 25
for epoch in range(num_epochs):
    # Set model to training mode and initialize running loss
    model.train()
    running_loss = 0.0

    # Loop over batches in the training dataset
    for batch in train_loader:
        inputs, labels = batch['spectrogram'].to(device), batch['label'].
        ↪to(device)

        # Zero the gradients
        optimizer.zero_grad()
        # Forward pass
        outputs = model(inputs)
        # Compute loss
        loss = criterion(outputs, labels)
        # Backward pass
        loss.backward()
        # Update parameters
        optimizer.step()
        # Accumulate the loss
        running_loss += loss.item()

    # Compute and log training loss
    training_loss = running_loss / len(train_loader)
    writer.add_scalar('Loss/train', training_loss, epoch)

    # Set model to evaluation mode and compute validation loss
    model.eval()
    val_loss = 0.0
    with torch.no_grad():
        # Loop over batches in the validation dataset
        for batch in val_loader:
            # Extract inputs and labels from the batch
            inputs, labels = batch['spectrogram'].to(device), batch['label'].
            ↪to(device)

            # Forward pass: compute model output
            outputs = model(inputs)
            # Compute loss
            loss = criterion(outputs, labels)
            # Accumulate the validation loss over all of the batches
            val_loss += loss.item()

    # Compute and log validation loss

```

```

validation_loss = val_loss / len(val_loader)
writer.add_scalar('Loss/validation', validation_loss, epoch)

# Check for early stopping
if validation_loss < best_val_loss:
    best_val_loss = validation_loss
    epochs_no_improve = 0

    # Save the model with the timestamp in the filename
    torch.save(model.state_dict(), f"{paths['models_dir']}/model_sch_lr_es/
↪model_best.pth")

else:
    epochs_no_improve += 1
    if epochs_no_improve >= patience:
        print(f'Early stopping triggered after {epoch + 1} epochs!')
        early_stop = True
        break

# Step the scheduler for learning rate adjustment
scheduler.step()

# Log training progress and learning rate
current_lr = scheduler.get_last_lr()[0]
writer.add_scalar('Learning Rate', current_lr, epoch)
print(f"Epoch {epoch+1}/{num_epochs}, Training Loss: {training_loss},
↪Validation Loss: {validation_loss}")
print("-" * 75)

# Check if training stopped early and close TensorBoard writer
if not early_stop:
    print(f"Training completed after {num_epochs} epochs.")
writer.close()

```

Epoch 1/25, Training Loss: 16.455200016498566, Validation Loss:
0.9361141140644367

Epoch 2/25, Training Loss: 0.7687073995669683, Validation Loss:
0.7818952111097482

Epoch 3/25, Training Loss: 0.5248074397444725, Validation Loss:
0.6259872294389285

Epoch 4/25, Training Loss: 0.271688986569643, Validation Loss:
0.6221039570294894

Epoch 5/25, Training Loss: 0.23109924035767715, Validation Loss:

0.6166144792850201

Epoch 6/25, Training Loss: 0.20983949775497118, Validation Loss:
0.6104520444686596

Epoch 7/25, Training Loss: 0.1814278454830249, Validation Loss:
0.609301372216298

Epoch 8/25, Training Loss: 0.17908216205736002, Validation Loss:
0.6131350443913386

Epoch 9/25, Training Loss: 0.1755180264512698, Validation Loss:
0.6150861153235803

Early stopping triggered after 10 epochs!